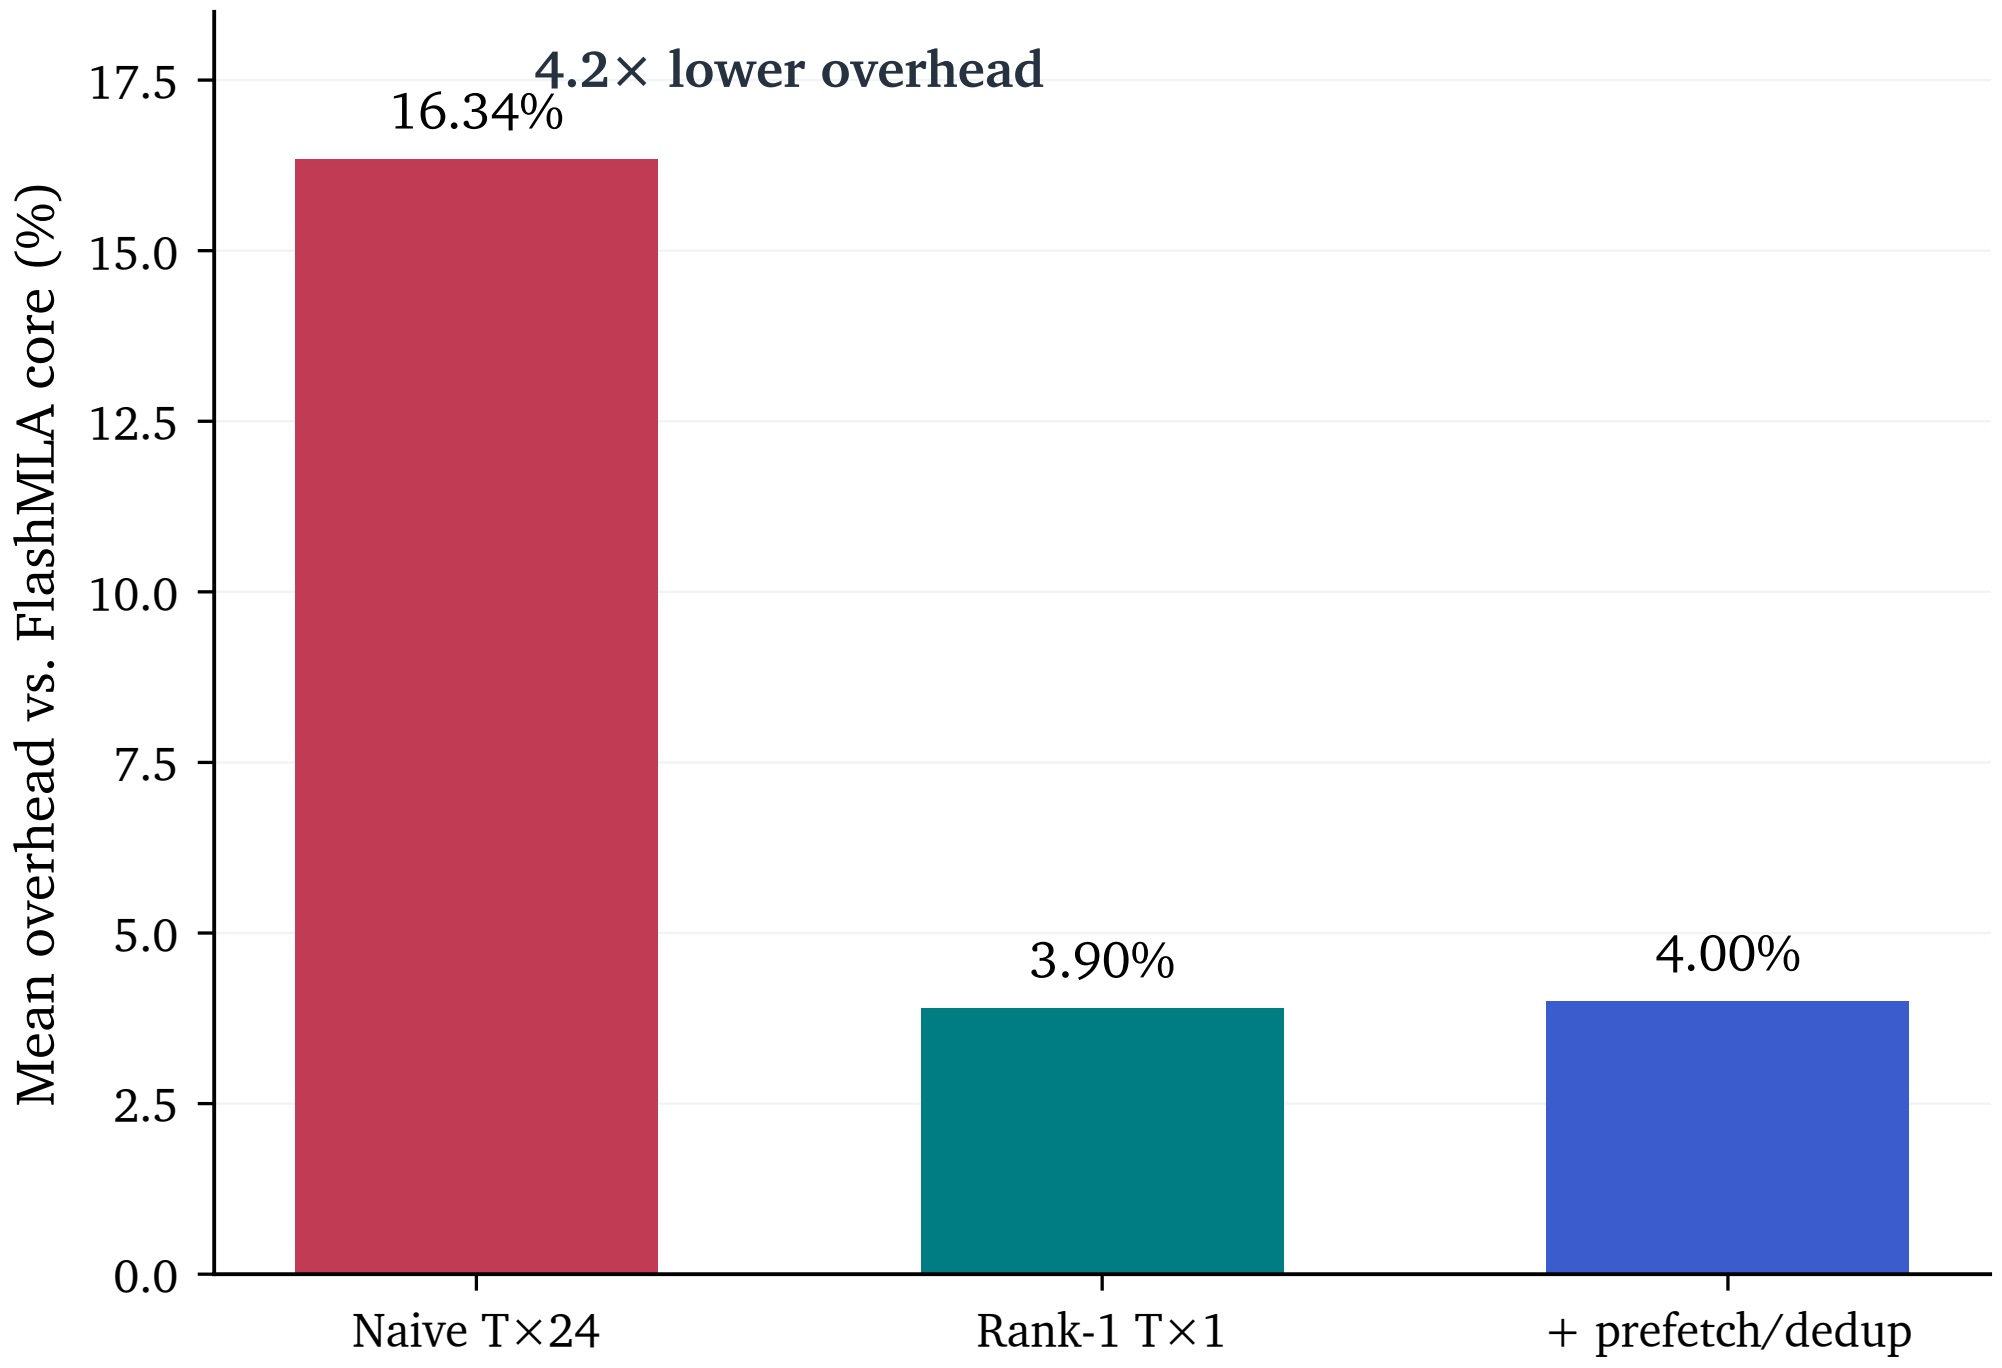
4.2× lower overhead
16.34%
3.90%
4.00%
Mean overhead vs. FlashMLA core (%)
17.5
15.0
12.5
10.0
7.5
5.0
2.5
0.0
Naive T×24
Rank-1 T×1
+ prefetch/dedup